

Project 1, John Clark, W205 Section 9 - Increasing Ridership at Lyft Bay Wheels

In this project, I will perform an analysis to determine how to increase ridership at Lyft Bay Wheels. We begin this effort by making a number of SQL queries and performing some visualizations.

Initialize Libraries

In [20]:

```
import pandas as pd
```

SQL Queries

We start out with SQL queries asking three questions:

- What is the size of this dataset?
- What is the earliest start date and time and latest end date and time for a trip?
- How many bikes are there:

The following queries will answer these questions

In [22]:

```
! bq query --use_legacy_sql=FALSE 'SELECT count(*) FROM `bigquery-public-data.san_francisco.bikeshare_trips`'

! bq query --use_legacy_sql=FALSE 'SELECT min(start_date), max(end_date), FROM `bigquery-public-data.san_francisco.bikeshare_trips`'

! bq query --use_legacy_sql=false 'select count(distinct(bike_number)) from `bigquery-public-data.san_francisco.bikeshare_trips`'
```

Waiting on bqjob_r59f98f92ef6d467b_00000179e7bf2e86_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 983648 |
+-----+
```

Waiting on bqjob_reb13804f737e5f5_00000179e7bf362d_1 ... (0s) Current status: DONE

```
+-----+-----+
| f0_ | f1_ |
+-----+-----+
| 2013-08-29 09:08:00 | 2016-08-31 23:48:00 |
+-----+-----+
```

Waiting on bqjob_r4dec8408492823d8_00000179e7bf3c02_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 700 |
+-----+
```

As we can see, the trips database has a total of 983,648 records in it. The earliest trip occurred on August 29th in 2013 and the latest trip occurred on August 31st of 2016. Given total unique bike numbers, there appear to be 700 bikes. Let's consider an additional couple of SQL queries to answer the following question:

- How many trips are in the morning vs. the afternoon?

Let us assume the morning is defined between 6:00AM and 11:59AM and the afternoon is defined as between 12:00pm and 4:59PM.

In [23]:

```
! bq query --use_legacy_sql=false 'SELECT count(*)from `bigquery-public-data.san_francisco.bikeshare_trips` where extract(hour from end_date AT TIME ZONE "UTC") between 6 AND 11'
! bq query --use_legacy_sql=false 'SELECT count(*) from `bigquery-public-data.san_francisco.bikeshare_trips` where extract(hour from end_date AT TIME ZONE "UTC") between 12 AND 16'
```

Waiting on bqjob_r2ff41df03b20782f_00000179e7bf79e9_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 387728 |
+-----+
```

Waiting on bqjob_r34d40c1bc1292087_00000179e7bf7fcd_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 250705 |
+-----+
```

Our query tells us that there are 387,728 trips in the morning and 250,705 trips in the afternoon.

Let us seek answers to some additional questions that tell us more about our dataset:

- Is there any revenue leakage due to bicycles being unavailable?
- What is the list of stations?
- What are the most popular trips?
- How many trips occur during peak rush hour times of 7:00PM to 9:59PM and 4:00PM to 6:59PM?
- How many trips make use of a subscription vs. non subscription?
- What is the distribution of trips across the different landmark locations (cities)?
- What days of the week do people most frequently use the bikes?
- Is there any seasonality present in ridership that may impact when offers should occur?

First, we look at the availability of bikes compared to the total number of bikes.

In [24]:

```
! bq query --use_legacy_sql=false 'select count(*) FROM `bigquery-public-data.san_francisco.bikeshare_status` where bikes_available = 0'
! bq query --use_legacy_sql=false 'select count(*) FROM `bigquery-public-data.san_francisco.bikeshare_status`'
```

Waiting on bqjob_r29a1102294cea88d_00000179e7bffc41e_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 850830 |
+-----+
```

Waiting on bqjob_r3097225181f84444_00000179e7bffc2e_1 ... (0s) Current status: DONE

```
+-----+
| f0_ |
+-----+
| 107501619 |
+-----+
```

We can see that 850,830 times, there have been zero bikes available at a station. However, when this is compared to the total number of status updates of 10,750,619, the number does not appear so large. In fact, there are no bikes available .0079% of the time or less than 1 percent.

Let's look next at the list of stations so we have an idea geographically as to where they are.

In [25]:

```
! bq query --use_legacy_sql=FALSE 'SELECT distinct station_id, name, landmark FROM `big
query-public-data.san_francisco.bikeshare_stations` ORDER BY landmark'
```

Waiting on bqjob_r5001f16f17247d2c_00000179e7c02ee1_1 ... (0s) Current status: DONE

station_id	name	landmark
32	Castro Street and El Camino Real	Mountain View
27	Mountain View City Hall	Mountain View
33	Charleston Park/ North Bayshore Area	Mountain View
30	Middlefield Light Rail Station	Mountain View
31	San Antonio Shopping Center	Mountain View
29	San Antonio Caltrain Station	Mountain View
28	Mountain View Caltrain Station	Mountain View
37	Cowper at University	Palo Alto
35	University and Emerson	Palo Alto
38	Park at Olive	Palo Alto
36	California Ave Caltrain Station	Palo Alto
34	Palo Alto Caltrain Station	Palo Alto
25	Stanford in Redwood City	Redwood City
21	Franklin at Maple	Redwood City
24	Redwood City Public Library	Redwood City
26	Redwood City Medical Center	Redwood City
23	San Mateo County Center	Redwood City
83	Mezes Park	Redwood City
22	Redwood City Caltrain Station	Redwood City
41	Clay at Battery	San Francisco
48	Embarcadero at Vallejo	San Francisco
46	Washington at Kearney	San Francisco
42	Davis at Jackson	San Francisco
45	Commercial at Montgomery	San Francisco
54	Embarcadero at Bryant	San Francisco
60	Embarcadero at Sansome	San Francisco
57	5th at Howard	San Francisco
65	Townsend at 7th	San Francisco
64	2nd at South Park	San Francisco
82	Broadway St at Battery St	San Francisco
73	Grant Avenue at Columbus Avenue	San Francisco
47	Post at Kearney	San Francisco
56	Beale at Market	San Francisco
51	Embarcadero at Folsom	San Francisco
49	Spear at Folsom	San Francisco
58	San Francisco City Hall	San Francisco
62	2nd at Folsom	San Francisco
63	Howard at 2nd	San Francisco
68	Yerba Buena Center of the Arts (3rd @ Howard)	San Francisco
66	South Van Ness at Market	San Francisco
70	San Francisco Caltrain (Townsend at 4th)	San Francisco
71	Powell at Post (Union Square)	San Francisco
75	Mechanics Plaza (Market at Battery)	San Francisco
76	Market at 4th	San Francisco
39	Powell Street BART	San Francisco
50	Harry Bridges Plaza (Ferry Building)	San Francisco
55	Temporary Transbay Terminal (Howard at Beale)	San Francisco
59	Golden Gate at Polk	San Francisco

72	Civic Center BART (7th at Market)	San Francisco
69	San Francisco Caltrain 2 (330 Townsend)	San Francisco
74	Steuart at Market	San Francisco
61	2nd at Townsend	San Francisco
67	Market at 10th	San Francisco
77	Market at Sansome	San Francisco
90	5th St at Folsom St	San Francisco
91	Cyril Magnin St at Ellis St	San Francisco
4	Santa Clara at Almaden	San Jose
84	Ryland Park	San Jose
8	San Salvador at 1st	San Jose
9	Japantown	San Jose
3	San Jose Civic Center	San Jose
13	St James Park	San Jose
10	San Jose City Hall	San Jose
16	SJSU - San Salvador at 9th	San Jose
7	Paseo de San Antonio	San Jose
6	San Pedro Square	San Jose
80	Santa Clara County Civic Center	San Jose
14	Arena Green / SAP Center	San Jose
5	Adobe on Almaden	San Jose
11	MLK Library	San Jose
12	SJSU 4th at San Carlos	San Jose
89	S. Market st at Park Ave	San Jose
88	5th S. at E. San Salvador St	San Jose
2	San Jose Diridon Caltrain Station	San Jose

A majority of the stations are in San Francisco and San Jose but we also see stations in Mountain View, Redwood City and Palo Alto. Notably, there are no stations in the East Bay.

Let's now consider the most popular trips.

In []:

```
! bq query --use_legacy_sql=FALSE 'SELECT end_station_name, start_station_name, landmark, count(*) as trip_freq FROM `bigquery-public-data.san_francisco.bikeshare_trips` a INNER JOIN `bigquery-public-data.san_francisco.bikeshare_stations` b ON a.end_station_id BETWEEN b.station_id AND b.station_id GROUP BY landmark, end_station_name, start_station_name ORDER BY trip_freq DESC'
```

We can see that 96 of the 100 most popular trips terminate in San Francisco. Two terminate in Mountain View and two terminate in San Jose.

Now let us look at how many trips occur at peak rush hour times of 7:00PM to 9:59PM and 4:00PM to 6:59PM.

In [4]:

```
%%bigquery

SELECT count(*) from bigquery-public-data.san_francisco.bikeshare_trips
where extract(hour from end_date AT TIME ZONE "UTC") between 7 AND 9

Query complete after 0.00s: 100% | 1/1 [00:00<00:00, 903.17query/s]
Downloading: 100% | 1/1 [00:01<00:00, 1.02s/rows]
```

Out[4]:

f0_

0 288598

In [5]:

```
%%bigquery
```

```
SELECT count(*) from `bigquery-public-data.san_francisco.bikeshare_trips`  
where extract(hour from start_date AT TIME ZONE "UTC") between 16 AND 18
```

```
Query complete after 0.00s: 100%|██████████| 2/2 [00:00<00:00, 483.35query/s]
```

```
Downloading: 100%|██████████| 1/1 [00:01<00:00, 1.23s/rows]
```

Out[5]:

	f0_
0	299626

There are 288,598 trips at the peak of morning rush hour and 299,626 during the peak of afternoon/early evening rush hour. These two numbers total to 60% of all the trips at Lyft Bay Wheels, therefore, targeting commuters appears to a good strategy of increasing sales volumes.

Let us now look at the breakdown of subscribers based on their subscription type.

In [30]:

```
! bq query --use_legacy_sql=false 'SELECT subscriber_type, count(subscriber_type) from  
`bigquery-public-data.san_francisco.bikeshare_trips` GROUP by subscriber_type'
```

```
Waiting on bqjob_r58e68c1cb16df251_00000179e7c4ef2d_1 ... (0s) Current status: DONE
```

subscriber_type	f0_
Customer	136809
Subscriber	846839

Based on this query, we can determine that 846,839 of the rides are by subscribers and only 136,809 are one time customers or those with a three day pass. Unfortunately this information is not very granular to inform us as to which types of promotions have been in place in the past, nor how effective those particular promotions have been.

What is the distribution of trips across the different landmark locations (cities)?

Next, we'll load the joined trips and stations table into a panda data frame and report off that. This next magic command saves it into a Pandas data frame and then we plot the results.

In [116]:

```
%%bigquery df
```

```
SELECT * FROM `bigquery-public-data.san_francisco.bikeshare_trips` a  
INNER JOIN `bigquery-public-data.san_francisco.bikeshare_stations` b  
ON a.end_station_id BETWEEN b.station_id AND b.station_id
```

```
Query complete after 0.00s: 100%|██████████| 1/1 [00:00<00:00, 941.91query/s]
```

```
Downloading: 100%|██████████| 983648/983648 [00:02<00:00, 466922.12rows/s]
```

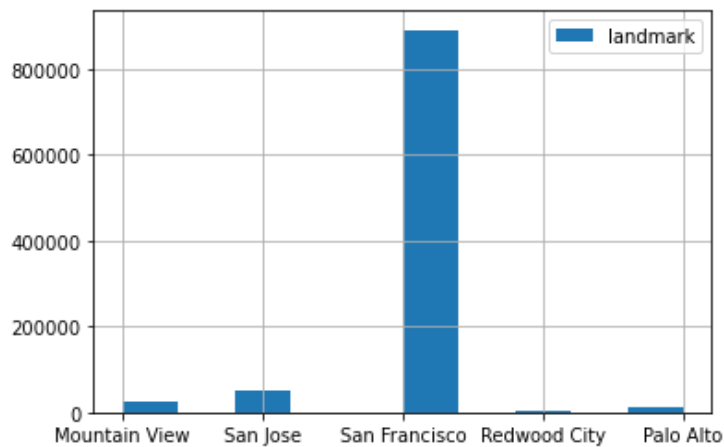
In [83]:

```
df.landmark.hist(legend = True)
```

Out[83]:

```
Out[83]:
```

```
<AxesSubplot:>
```



That is an enormous amount of trips centered around San Francisco relative to other locations.

Which days of the week are most popular?


```
In [117]:
```

```
# Compute day of week and add to dataframe
df['day_of_week'] = df['start_date'].dt.day_name()

#Show histogram demonstrating most popular days
df.day_of_week.hist(bins = 7, figsize = [8.8, 4.8])
```

```
Out[117]:
```

```
<AxesSubplot:>
```



The weekdays appear to be the most popular by a large margin, consistent with a theory that bike rides are occurring during the week and for commuting.

Weekday Trips by Month

```
In [128]:
```

```
#Sort values by day of week and start date
```

```
df.sort_values(["day_of_week", "start_date"])

#Drop weekends
indexNames = df[ (df['day_of_week'] == 'Saturday') | (df['day_of_week'] == 'Sunday') ].index
df.drop(indexNames , inplace=True)

len(df)
```

Out[128]:


871994

In [129]:

```
df.hist(column='start_date', bins = 36, figsize = [11,4.8])
```

Out[129]:

array([[<AxesSubplot:title={'center':'start_date'}>]], dtype=object)



Since we are focused on commuter offers, the weekends are removed from the dataframe and the resulting histogram of the trips shows a definite pattern of seasonality. The least popular months are in the middle of winter. Volume of rides appears to climb through the fall timeframe. This is likely due to the weather being colder and possibly additional rain during the winter.

Conclusion and Recommendations

Conclusions

It is clear from the data that an opportunity exists to target customers who are commuting to and from their jobs. There is a seasonality present in the ridership with certain months during the winter appearing as "down" months. In this analysis, we have identified the most popular trips, the busiest stations, the most important commuting windows, with weekdays being a priority and a high concentration of trips overall in San Francisco. Opportunities may exist to broaden the number of stations to certain locations that may be underserved inside or outside of San Francisco. There isn't sufficient granularity in the subscription data within the dataset and the lack of revenue data makes it somewhat difficult to evaluate which specific offers may be the most effective. Station real estate may be one of the most important assets currently held by the company and for that reason, I provide several recommendations related to those locations.

Recommendations

Expand the use of the subscriber_type field to include more granular identification of the types of promotional campaigns, if any, that the rider is subscribing to. It is difficult to determine what promotional campaigns are working or not at the present moment because the data isn't in the dataset to evaluate which types of campaigns are providing the greatest return. It's further possible that a subscriber should participate in more than one program at the same time, which isn't possible to tell at the moment. Tracking this kind of information would provide more granular information on what campaigns are working and when.

Correlate revenue with trip data in order to enable stronger patterns to be discerned on potential opportunities. Because this is a public dataset, it isn't surprising that the company did not provide revenue information. This information, if available could be used to correlate data from trips as well as the various promotional programs in order to best determine what types of offers may optimally increase revenue and improve customer service.

Time promotional offers to occur in early spring. Promotional offers should likely be timed to begin in early March as this is when the weather becomes more amenable to bicycle riding. Increasing ridership at this stage of the year may likely have benefits through the fall.

Target mobile ads to individuals close to Lyft bike stations. We have geolocation information available to us for each of the stations. In addition, mobile ad targeting can also identify individuals who live close to those stations. It is recommended to correlate these two sets of information and partnering with Google and with other data aggregators or advertisers, to structure an advertisement campaign targeted at these individuals.

Expand into other markets outside San Francisco. Based on the statistics, it seems apparent that there may be a large number of additional opportunities outside of the San Francisco area that may provide an opportunity for additional ridership, especially including locations near the Bay Area Rapid Transit (BART) train stations in the East Bay, areas around Oakland and surrounding communities, and other areas with a high concentration of employers. Riders in these communities could be enticed through bike stations close to or colocated with Bart Stations and advertising done on BART Trains. BART offers a number of advertising opportunities:

<https://www.bart.gov/about/business/advertising>.

Consider expanding stations to ferry locations in the Bay Area. To avoid the gridlock, Bay Area residents traveling into San Francisco are increasingly using ferries. In order of priority, I would consider adding stations at Larkspur (3,048,733 trips in 2019), Alameda/Oakland (1,384,300 trips in 2019) and Vallejo (1,081,665 trips in 2019).

Launch a campaign targeted at increasing the number of corporate memberships with locations that are near bike stations and BART stations. Corporate memberships provide an attractive opportunity for employers to offer a benefit to their employees and support the health and well being of their employees. For Lyft, they offer an incremental source of income. Large campus sites may also offer a potential opportunity for additional stations, making it more convenient for the employees.